

Introduction to neural networks

Training aspects

Juan Irving Vásquez
jivg.org

Centro de Innovación y Desarrollo Tecnológico en Cómputo
Instituto Politécnico Nacional

March 19, 2026©



Contents

- 1 Performance measurement
- 2 Dataset split
- 3 Performance metrics
- 4 Overfitting
- 5 Regularization

Roadmap

- 1 Performance measurement
- 2 Dataset split
- 3 Performance metrics
- 4 Overfitting
- 5 Regularization

Performance measurement

- The central challenge in machine learning is that we must perform well on new, **previously unseen** inputs.



- The ability to perform well on previously unobserved inputs is called **generalization**.
 - Dataset split
 - Metric
 - Significance

¹ <https://www.bbc.com/news/world-us-canada-43604440>

² <https://www.formacar.com/en/tesla-63631-article>

Performance measurement

We will introduce a validation step and a test in order to measure generalization.




Figure: Train and test workflow

Roadmap

- 1 Performance measurement
- 2 Dataset split**
- 3 Performance metrics
- 4 Overfitting
- 5 Regularization

Dataset split




Figure: Dataset split [HJC]

Training and validation set split

N-fold validation

- Lets suppose a dataset, $X_{\Omega} = \{x_1, \dots, x_m\}$, with positive examples
- We could use a subset of our set of observations as an estimate of the performance

$$X = x_1, \dots, x_n$$

$$X' = x_{n+1}, \dots, x_m$$

, X is the **training set** and X' is usually called the **validation set** [Smola].

- Typically, one uses about 80% of the training data for training and 20% for validation [GoodFellow].

Roadmap

- 1 Performance measurement
- 2 Dataset split
- 3 Performance metrics**
- 4 Overfitting
- 5 Regularization

Common measurements for categorical outputs

- Accuracy
- Precision
- Recall
- F1 score

Confusion matrix

- Computation:

Ground truth	Positive Prediction	Negative Prediction
Positive Class	TP	FN
Negative Class	FP	TN

Table: Confusion matrix

- $TP + FN$ sums 100 % of the positive class
- $FP + TN$ sums 100 % of the negative class

Example

Animal	 Barn owl  Chihuahua  Sheep dog
Not animal	 Apple  Muffin  Mop

Figure: Dataset of animals and no animals. ³

³<https://www.labelf.ai/blog/what-is-accuracy-precision-recall-and-f1-score>

Prediction concepts

For binary classes⁴

Network input	Prediction	Concept
	'animal'	true positive
	'not animal'	false negative
	'animal'	false positive
	'not animal'	true negative

⁴Consider the class 'animal' as the positive class.

Confusion matrix

		Predicted			
		Animal	Not animal		
Actual	Animal			True Positives	
	Not animal			True Negatives	
				False Positives	
				False Negatives	

Figure: Confusion matrix example.

Confusion matrix

		Predicted			
		Animal	Not animal		
Actual	Animal			True Positives	2
	Not animal			True Negatives	2
				False Positives	1
				False Negatives	1

Figure: Confusion matrix example.

Performance metrics

Accuracy

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN} \quad (1)$$

TP represents true positives, TN true negatives, FP false positives, and FN false negatives.

Performance metrics

Precision

$$\text{Precision} = \frac{TP}{TP + FP} \quad (2)$$

In this formula, TP represents true positives and FP false positives. This metric measures the proportion of positive identifications that were actually correct.

Performance metrics

Recall

$$\text{Recall} = \frac{TP}{TP + FN} \quad (3)$$

TP represents true positives and FN false negatives. Recall measures the model's ability to find all relevant positive instances in the dataset.

Performance metrics

F1-score

The F1-score, which is the harmonic mean of precision and recall, can be expressed as:

$$F1 = 2 \times \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}} \quad (4)$$

Thus, the F1-score combines both precision and recall into a single metric that accounts for both false positives and false negatives. It is especially useful in situations where the balance between precision and recall is critical.

Significance





Definition

A result is statistically significant if the probability of it occurring by chance is very low, typically defined by a threshold α (e.g., 0.05).

Core Concepts:

- **Null Hypothesis (H_0):** No real difference between models.
- **p-value:** $P(\text{Data}|H_0)$. If $p < \alpha$, we reject H_0 .

Application in ML (Model Comparison): To compare error rates ϵ_1 and ϵ_2 using a Z-test:

$$Z = \frac{|\hat{\epsilon}_1 - \hat{\epsilon}_2|}{\sqrt{\frac{\hat{\epsilon}(1-\hat{\epsilon})}{n_1} + \frac{\hat{\epsilon}(1-\hat{\epsilon})}{n_2}}} \quad (5)$$



Ensures that performance gains are robust and reproducible.

Roadmap

- 1 Performance measurement
- 2 Dataset split
- 3 Performance metrics
- 4 Overfitting**
- 5 Regularization


Overfitting problem

Which model is better?





Overfitting problem

The importance of testing




Overfitting problem

The importance of testing



Overfitting problem

The importance of testing



Overfitting and generalization

- **Underfitting** occurs when the model is not able to obtain a sufficiently low error value on the training set.
- **Overfitting** occurs when the gap between the training error and test error is too large.
- A model's **capacity** is its ability to fit a wide variety of functions.

Regularization

Identify overfitting




Figure: Loss as epochs advance.

Roadmap

- 1 Performance measurement
- 2 Dataset split
- 3 Performance metrics
- 4 Overfitting
- 5 Regularization**

- **Regularization** is any modification we make to a learning algorithm that is intended to reduce its generalization error but not its training error.
 - Early termination
 - L1, L2
 - Dropout

Regularization

Early termination




Figure: Performance as epochs advance.

L2 Regularization (Weight Decay)

Adds a penalty equal to the square of the magnitude of coefficients.

The Math:

$$\mathcal{L}(W) = E(W) + \lambda \sum_{j=1}^n w_j^2$$

Key Insights:

- Penalizes large weights.
- Leads to **small, non-zero** weights.

L1 Regularization (Lasso)

Adds a penalty equal to the absolute value of the magnitude of coefficients.

The Math:

$$\mathcal{L}(W) = E(W) + \lambda \sum_{j=1}^n |w_j|$$

Key Insights:


- Encourages **sparsity** (many weights become exactly zero).
- Useful for **feature selection**.

A stochastic technique mainly used in Deep Learning.


- **How it works:** During training, randomly "deactivate" a percentage p of neurons in each forward pass ⁵.
- **Why it works:**
 - Prevents **co-adaptation** of neurons.
 - Forces the network to learn redundant representations.
 - Acts as an approximate model ensemble (averaging many sub-networks).

⁵All neurons are used during inference, but weights are scaled by p .

Standard Network



Network with Dropout



Model capacity selection





Figure: Selecting depth

Model capacity selection

- Usually select the simplest model
- The optimal model is the one that has the same capacity than the problem.



- Chapter 5 Machine Learning Basics. Ian Goodfellow, et al. Deep Learning. MIT.
- Chapter 2 Training feed-forward neural networks. Buduma. Fundamentals of deep learning.

References

- Goodfellow, I., Bengio, Y., Courville, A., & Bengio, Y. (2016). Deep learning (Vol. 1). Cambridge: MIT press.
- Sucar, Luis Enrique. Probabilistic Graphical Models: Principles and Applications. Springer, 2015.
- Alex Smola and S.V.N. Vishwanathan, Introduction to Machine Learning, Cambridge University Press.
- Self Driving Car Nanodegree
- Srivastava, N., Hinton, G., Krizhevsky, A., Sutskever, I., & Salakhutdinov, R. (2014). Dropout: a simple way to prevent neural networks from overfitting. The Journal of Machine Learning Research, 15(1), 1929-1958.
- <https://houstonjobconnection.com/article/what-is-a-dataset-in-machine-learning>